# Fronthaul Load Management in Open RAN

Part A: shrinking the load with matrix-decomposition compression
Part B: sharing one fronthaul link between cells with forecast-aware budgets

B.Tech. Project — mid-term evaluation

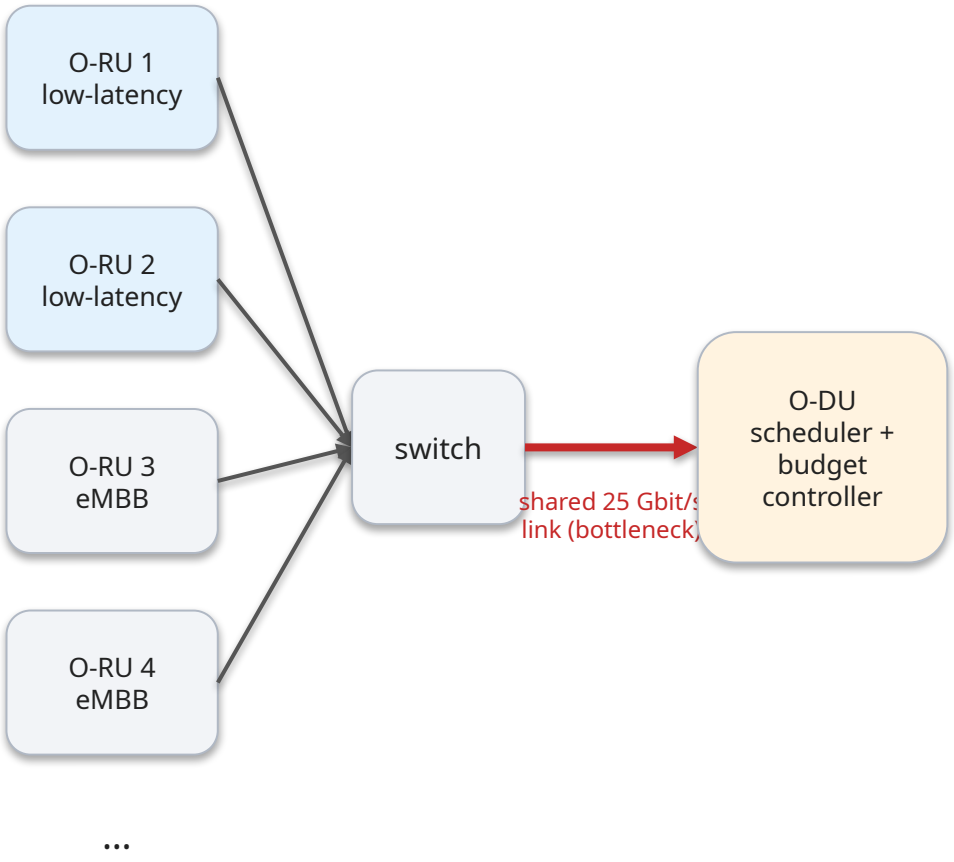
Student: <name> | Supervisor: <name> | Department of <dept>

*What exists today: compression study + matrix-maths benchmark (Part A); working simulator, 7 controllers, trained forecaster, 300+ experiment runs, tests and this deck (Part B). All generated from the repository.*

# The problem: the fronthaul carries a lot of data



- In the 7-2x split, the radio unit (O-RU) sends radio samples (IQ data) to the processing unit (O-DU)
  - Every scheduled resource block costs about 3,387 bits on the link (9-bit compression, 8% packet overhead)
  - One 100 MHz cell with 4 antenna layers can need up to 7.4 Gbit/s
- Several cells share one link that is smaller than their total peak — on purpose (statistical multiplexing)
  - Our setup: 8 cells on a 25 Gbit/s link; total peak = 2.4 × the link
- When bursts coincide, some data must wait or be dropped. Low-latency cells can wait ~2 ms at most
- **Two ways to manage the load:**
  - A. Make the data smaller — compression (matrix decomposition, quantisation)
  - B. Share the link well — decide every few milliseconds how much each cell may send



# Project overview: two parts, two levers



## Part A — make the data smaller (earlier work)

- Compression of the IQ matrix Y (64 antennas × 1200 subcarriers) per slot
- Baselines: O-RAN block floating point (BFP), truncated SVD
- Studied: RAS-BFP (random sketch + QR → low rank), CSEE (keep the strongest delay taps)
- ACAFS: pick the functional split per user from the channel rank
- Benchmark of the matrix maths behind it: LU, QR, SVD, Newton-Schulz, learned inverters (n = 10 … 500)
- Output: 4 figure suites, benchmark CSV, report. Compression source package is not in this repo (see status)

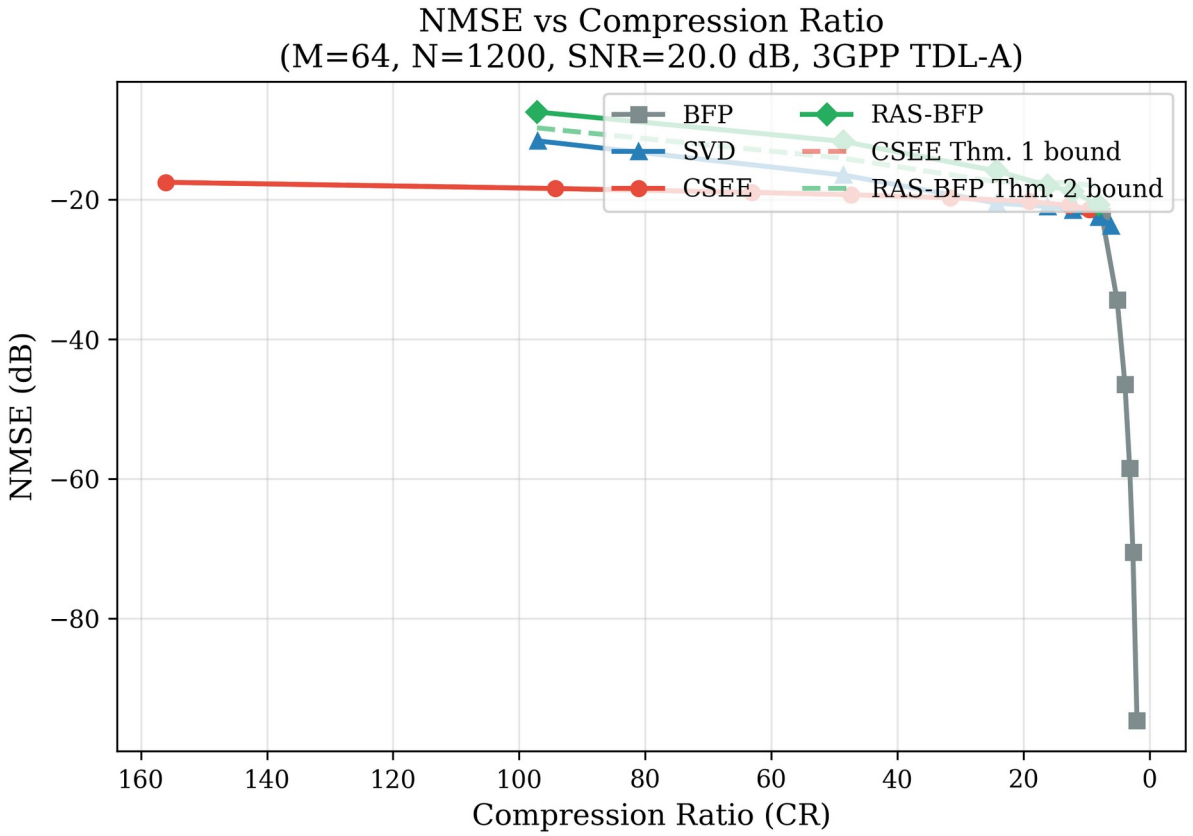
## Part B — share the link well (this term)

- Slot-level simulator of 8 cells on one 25 Gbit/s link, with queues, deadlines and priorities
- Every 10 ms a controller sets how many bits per slot each cell may send
- 5 reference controllers, from 'equal share' to 'knows the future'
- Proposed: forecast each cell's need as a range (quantiles), then split the link so that every cell is cut at the same weighted risk
- 200 test runs on 5 traffic scenarios, ablations, 14 unit tests
- Output: raw metrics, figures, this deck — all regenerated by scripts

How they connect: compression decides how much load reaches the link (Part A lever). Allocation decides how the loss is shared when bursts coincide (Part B lever). Slide 15 shows both levers on the same traffic.

# Part A: compressing IQ data with matrix decomposition

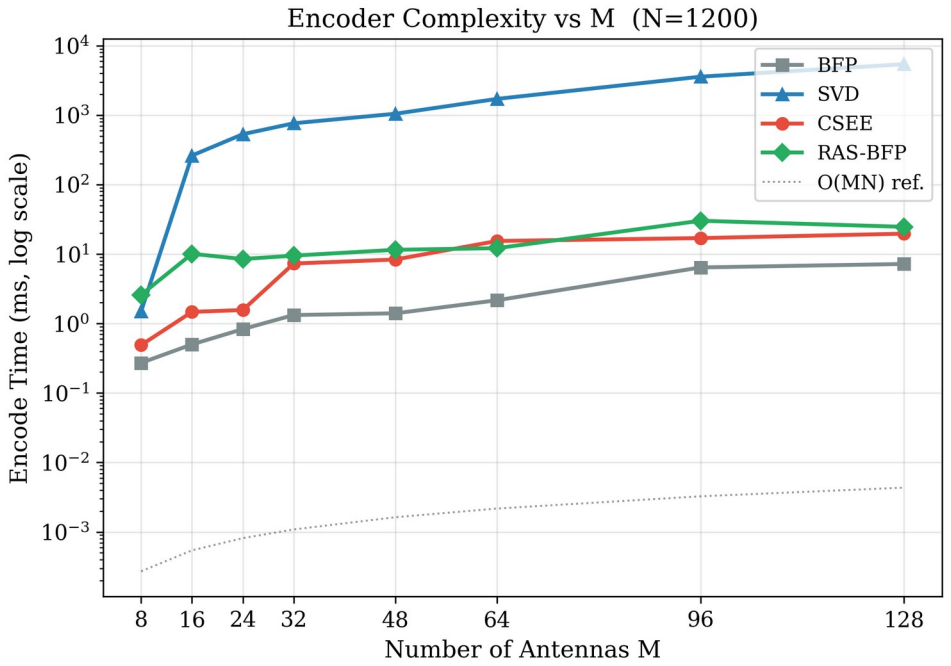





- Setup: Y ∈ C^(64 × 1200) per slot, 3GPP TDL-A channel, SNR 20 dB (simulated)
- BFP (O-RAN standard): shared exponent + short mantissa. ~4× smaller, −46 dB error, ~3 ms
- Truncated SVD: keep the r strongest directions. ~16× at −21 dB, but 18-30 ms per slot
- RAS-BFP: random sketch Ω·Y → thin QR → Y ≈ L·$Q^H$, then BFP. ~16× at −18 dB, ~7 ms (3-4× faster than SVD)
- CSEE: IFFT along subcarriers → keep top-K delay taps → BFP. ~13× at −21 dB (~5 ms); ~63× at −19 dB with K = 60
- ACAFS: choose split 7.2x / 7.1 / 6 per user from the channel rank → in simulation 0 % saving below ~15 dB SNR, 55-75 % above (vs fixed split 6)
- **Status: figures and report exist (results/figures, docs/mid_evaluation_report.md); the src/ package is missing from the repo, so these numbers are quoted, not re-run**

# Part A: how fast is the matrix maths? (measured on CPU)





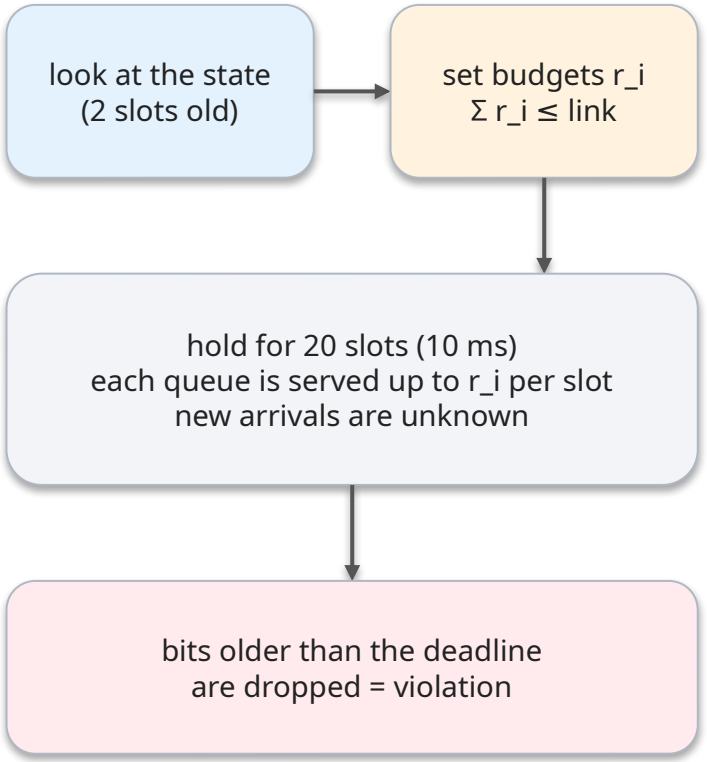

| Method | n=100: median ms | n=100: rel. error | n=500: median ms | n=500: rel. error |
|---|---|---|---|---|
| LU via LAPACK (numpy.linalg.inv) | 0.17 | 5e-08 | 18.00 | 5e-08 |
| LU (scipy) | 2.87 | 1e-07 | 56.29 | 2e-07 |
| QR | 3.94 | 3e-07 | 102.23 | 3e-07 |
| SVD | 2.76 | 1e-07 | 106.92 | 2e-07 |
| Newton-Schulz iteration | 1.01 | 2e-07 | 70.30 | 3e-07 |
| Learned iteration (InverseNet-Ultra) | 0.46 | 1e-06 | 33.79 | 1e-06 |
| Direct neural net (InverseNet-MLP) | 4.58 | 5e-01 | n/a | n/a |

- Encoder time grows fastest for SVD; sketch (RAS-BFP) and delay-domain (CSEE) encoders stay 10-100× below it
- Inversion benchmark (right, medians, 200 held-out matrices per size, well-conditioned): LU via LAPACK is fastest and exact
- A direct neural net that outputs the inverse does not work (≈ 50 % error). A learned Newton-Schulz iteration does (≈ 1e-6 at n = 100) and has a fixed, predictable cost
- Lesson for slot deadlines (0.5 ms): n ≈ 100 matrices fit per slot on a CPU; n ≈ 500 do not → do them less often, incrementally, or on an accelerator
- Lesson for Part B: predictable compute time matters as much as average speed — the allocation rule below is a closed form plus one scalar search

# Part B: what is decided, when, and with what information



- Time: slots of 0.5 ms. Data: bits. Link: bits per slot
- 8 cells: 3 low-latency (must be served within 4 slots = 2 ms, priority 3) and 5 eMBB (20 slots = 10 ms, priority 1)
- Fronthaul bits = user bits ÷ spectral efficiency × bits per resource block — so fronthaul load ≠ user traffic
- Each cell has a first-in-first-out queue in the O-DU; bits older than the deadline are dropped and counted
- Link rule: budgets add up to at most 97 % of 25 Gbit/s (3 % kept for control traffic); no cell gets more than its own peak
- Every T = 20 slots (10 ms) the controller sets the budgets; they are fixed for the interval
- The controller sees the state 2 slots late; it never sees future arrivals (checked by a unit test)
- **Why decide in advance? A budget is a reservation held for 10 ms; the traffic in that window is not known yet**




*Sources: O-RAN WG4 CUS-plane spec (7-2x, BFP, timing windows); Larsen et al., IEEE COMST 2019 (bit rate depends on load, ~8 % Ethernet overhead); Pérez et al. 2018/2019 (packet-switched aggregation).*

# What others have done, and the gap



- **A. The fronthaul itself (standards and surveys)**
  - O-RAN WG4 CUS-plane spec (7-2x, BFP, timing windows) • Larsen et al., IEEE COMST 2019: bit rate varies with load
  - Pérez et al., JOCN 2018 / IEEE Access 2019 (packet aggregation) • Wang & Zhou, IEEE Commun. Lett. 2017 (multiplexing gain)
  - Lagén et al., IEEE Commun. Mag. 2022: several 7-2x cells share one link, compression is adjusted per slot — reactive, no forecast
- **B. Forecast first, then allocate capacity**
  - Bega et al., DeepCog (INFOCOM 2019 / JSAC 2020), AZTEC (INFOCOM 2020): deep-learning forecasts for network slices
  - Predictive bandwidth allocation for PON fronthaul (Mikaeil et al. 2018; Zhang et al. 2019); Kavehmadavani et al., TWC 2023
  - Common pattern: one predicted number + a fixed safety margin; no per-cell deadlines
- **C. Forecasts with uncertainty**
  - Cohen et al., IEEE WCL 2023 (conformal prediction for one URLLC stream) • Gibbs & Candès, NeurIPS 2021 (adaptive conformal)
  - Kasuluru et al. 2024 [preprint]: probabilistic PRB-load forecasts; the percentile is picked globally
- **Gap: nobody turns a forecast RANGE of deadline-sensitive demand into per-cell budgets under one shared link, and checks it against strong reactive and point-forecast controllers**

# What we measure, and what 'demand' means here



- Decision every 10 ms: budgets $r_i \geq 0$, $\Sigma r_i \leq$ link capacity, $r_i \leq$ cell peak
- **Main score (fixed before tuning): priority-weighted deadline-violation ratio**
  - = (3 × dropped low-latency bits + dropped eMBB bits) / (3 × arrived low-latency bits + arrived eMBB bits)
  - Also reported: violations per class, link utilisation, queueing delay, fairness (Jain index), decision time
- **'Demand' of a cell = deadline-feasible rate r*: the smallest constant rate that serves the next arrivals with no bit waiting longer than its deadline D**
  - r* = max over time windows of (bits in window) / (window length + D) — a standard network-calculus bound
  - For eMBB (D = 10 ms) r* ≈ average rate. For low-latency cells (D = 2 ms) r* ≈ peak 2-ms burst rate → hard to predict
  - Budget needed = rate to clear the current backlog in time + r* of the arrivals we have not seen yet
- Why a single predicted number is not enough: r* is skewed (usually low, sometimes very high). The median under-reserves; a fixed margin wastes capacity on every cell

# Reference controllers (all implemented and tuned on validation)



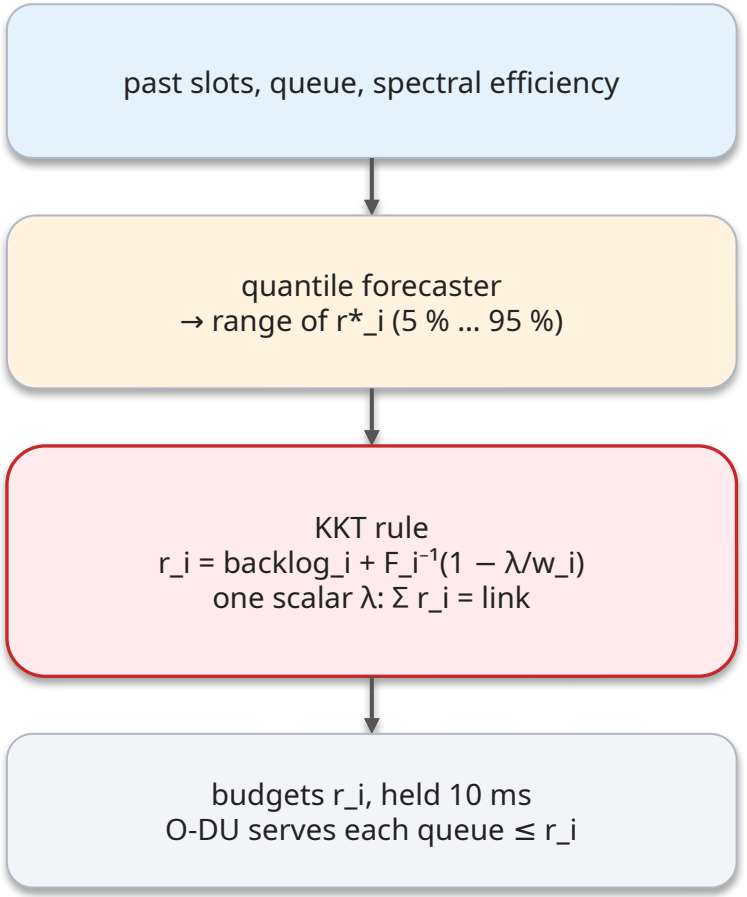
| Controller | How it sets budgets | What it uses | Tuned |
|---|---|---|---|
| Static equal share | link ÷ 8 per cell (spare capacity redistributed) | nothing | — |
| Reactive proportional | backlog / T + running average of demand; proportional split; spare redistributed | history | margin, window |
| MaxWeight (queue-weighted) | greedy fill in order of priority × queue length (optional head-of-line-age urgency) | queues, priorities | margin, HOL weight |
| Deadline-aware reactive (strong) | backlog-clearing rate + r* of the last interval; priority weights; water-fill | history, deadlines, priorities | margin, #intervals |
| Point forecast + water-fill (ablation) | same as above, but r* comes from the forecast MEDIAN | + trained model | margin |
| Fixed global quantile + water-fill | same as above, but r* = one tuned forecast percentile, identical for every cell | + trained model | level |
| Oracle water-fill (reference) | true r* of the next interval, same water-fill. Perfect point forecast; not a bound on other rules | future arrivals | — |

- **Tuning on validation seeds 2000-2003 (scenarios moderate/high); test seeds were never used:**
  - reactive_prop → {'margin': 0.0, 'ewma_window': 20} | queue_aware → {'margin': 0.0, 'num_horizons': 1} | point_forecast → {'margin': 0.0} | proposed_cal → {'gamma': 0.005}
  - Finding: every safety margin above 0 was worse. Inflating a single number hurts when the shortage is split proportionally — the cells that over-claim get more
- Same traffic traces, same delayed observations, spare capacity always redistributed, for every controller

# Proposed method: forecast a range, cut every cell at equal risk



- **1. Forecaster: six quantiles (5, 25, 50, 75, 90, 95 %) of each cell's r* for the next interval**
  - – One gradient-boosted tree model per quantile; 28 features from past slots only; trained on training seeds only
  - – Targets are divided by the cell's fixed peak, so nothing is fitted on test data
- **2. Allocation rule: maximise the weighted expected served demand under the forecast**
  - – max $\Sigma_i w_i \cdot E[\min(\text{backlog}_i + r^*_i, r_i)]$ subject to $\Sigma r_i \le \text{link}$, $0 \le r_i \le \text{peak}_i$
  - – Solution (KKT): $r_i = \text{backlog}_i + F_i^{-1}(1 - \lambda / w_i)$, clipped to $[0, \text{peak}_i]$; $\lambda$ found by bisection so that $\Sigma r_i = \text{link}$
  - – In words: every cell is cut at the same weighted risk of needing one more bit
  - – → wide forecast or high priority ⇒ more margin; confident forecast ⇒ ≈ median; backlog is always funded first
- Optional: an online calibration step (ACI) that shifts a cell's quantile level after misses — tested as an ablation
- Cost: a few microseconds for 8 cells; budgets always feasible (proofs in docs/fhlm/method.md)




*Borrowed: quantile boosting, ACI, water-filling.*
*Our part: r* as the forecast target + the KKT rule over the forecast range under the link constraint, tested against strong baselines.*

# Experimental setup (synthetic traffic, clearly labelled)

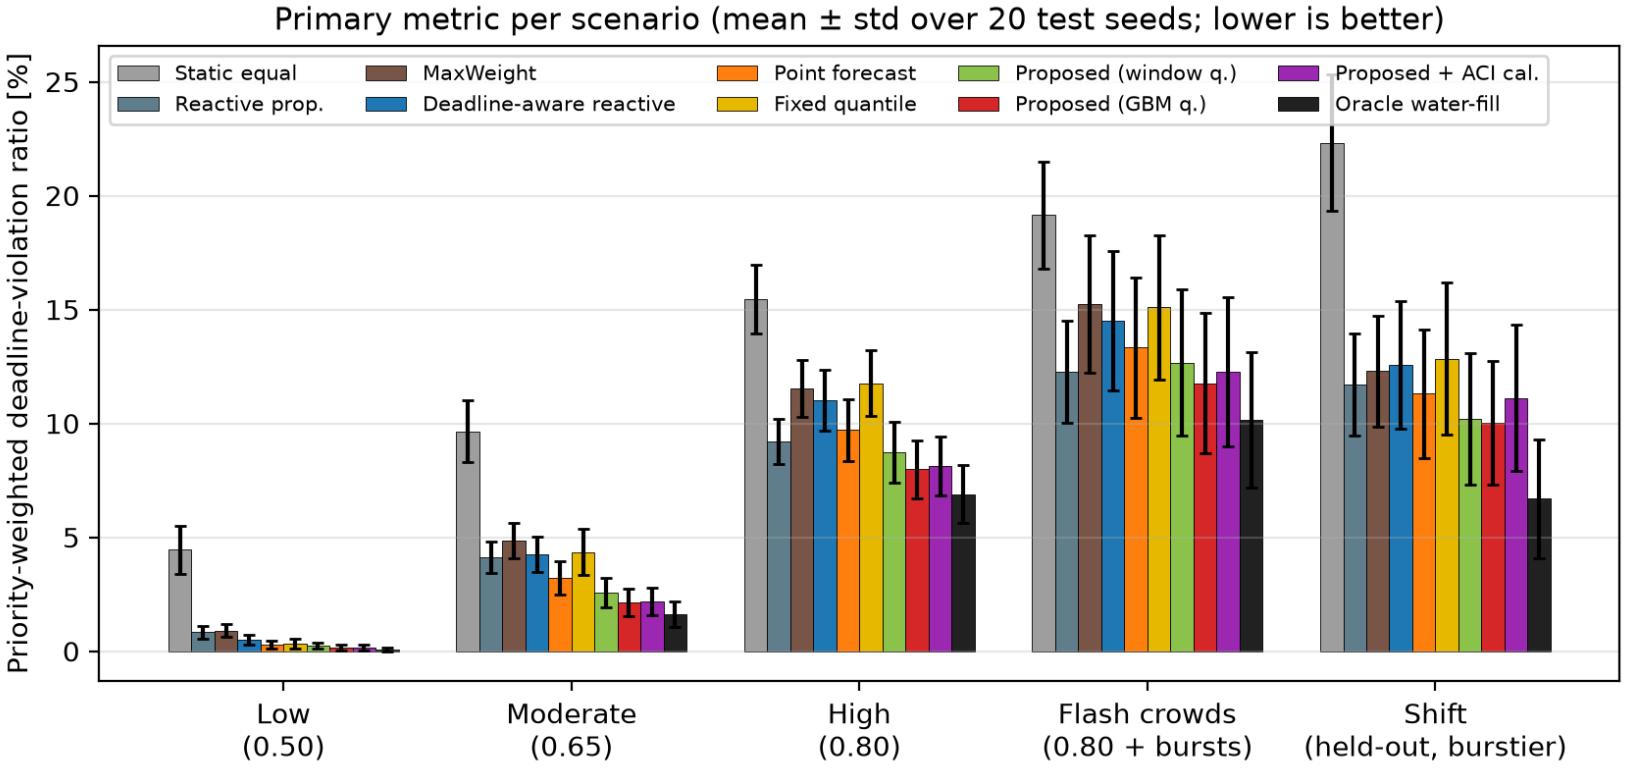


| Scenario | Mean load / link | Traffic | Used for tuning? |
|---|---|---|---|
| low | 0.50 | normal | no |
| moderate | 0.65 | normal | yes (validation seeds) |
| high | 0.80 | normal | yes (validation seeds) |
| overload_bursts | 0.80 | flash crowds: load ×1.8 for about 15 % of the time (more than the link can carry) | no |
| shift (held-out) | 0.80 | bursts 1.5× stronger and 1.5× longer, heavier tail — never seen in training or tuning | no |

- Traffic (synthetic): smooth Gamma part + heavy-tailed ON/OFF bursts + slow load changes (~1 s) + spectral-efficiency drift
- Seeds: training 1000-1015 (loads 0.5-0.95) | validation 2000-2003 | test 3000-3019 — no overlap
- Each run: 20,000 slots = 10 s after a 1,000-slot warm-up; T = 20 slots; 10 controllers × 5 scenarios × 20 seeds = 1000 runs
- Same traffic and same delayed observations for every controller; the forecaster is frozen before testing
- Absolute violation ratios are high on purpose (bursty traffic at 0.8 load); the paired differences between controllers are the result

# Results: main score per scenario (20 test seeds, mean ± std)

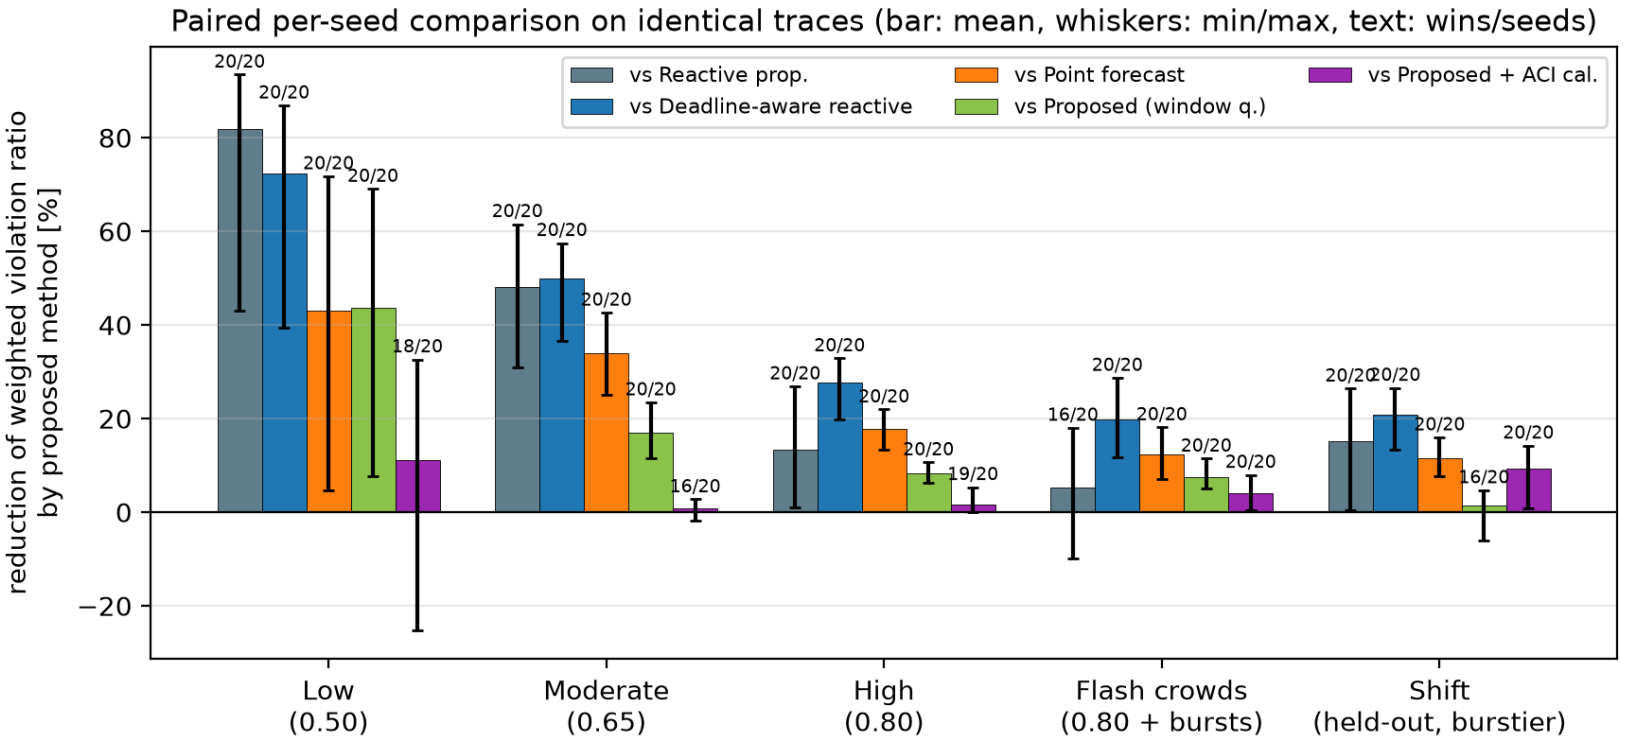





- **Proposed = best real controller in 5 of 5 scenarios**
  - vs point forecast (same model, median only): −18 % (high), −11 % (shift)
  - vs deadline-aware reactive: −27 % (high), −20 % (shift)
  - vs reactive proportional: −13 % (high) but only −4 % under flash crowds (16/20 seeds, p = 0.03) → small gain
  - vs MaxWeight: −31 % (high); vs fixed global quantile: −32 % (high) — per-cell levels beat one global level
- Gap to the oracle at high load: 8.0% vs 6.9%
- Error bars are wide because burst timing differs by seed → compare seed by seed (next slide)

# Results: seed-by-seed comparison on the same traffic

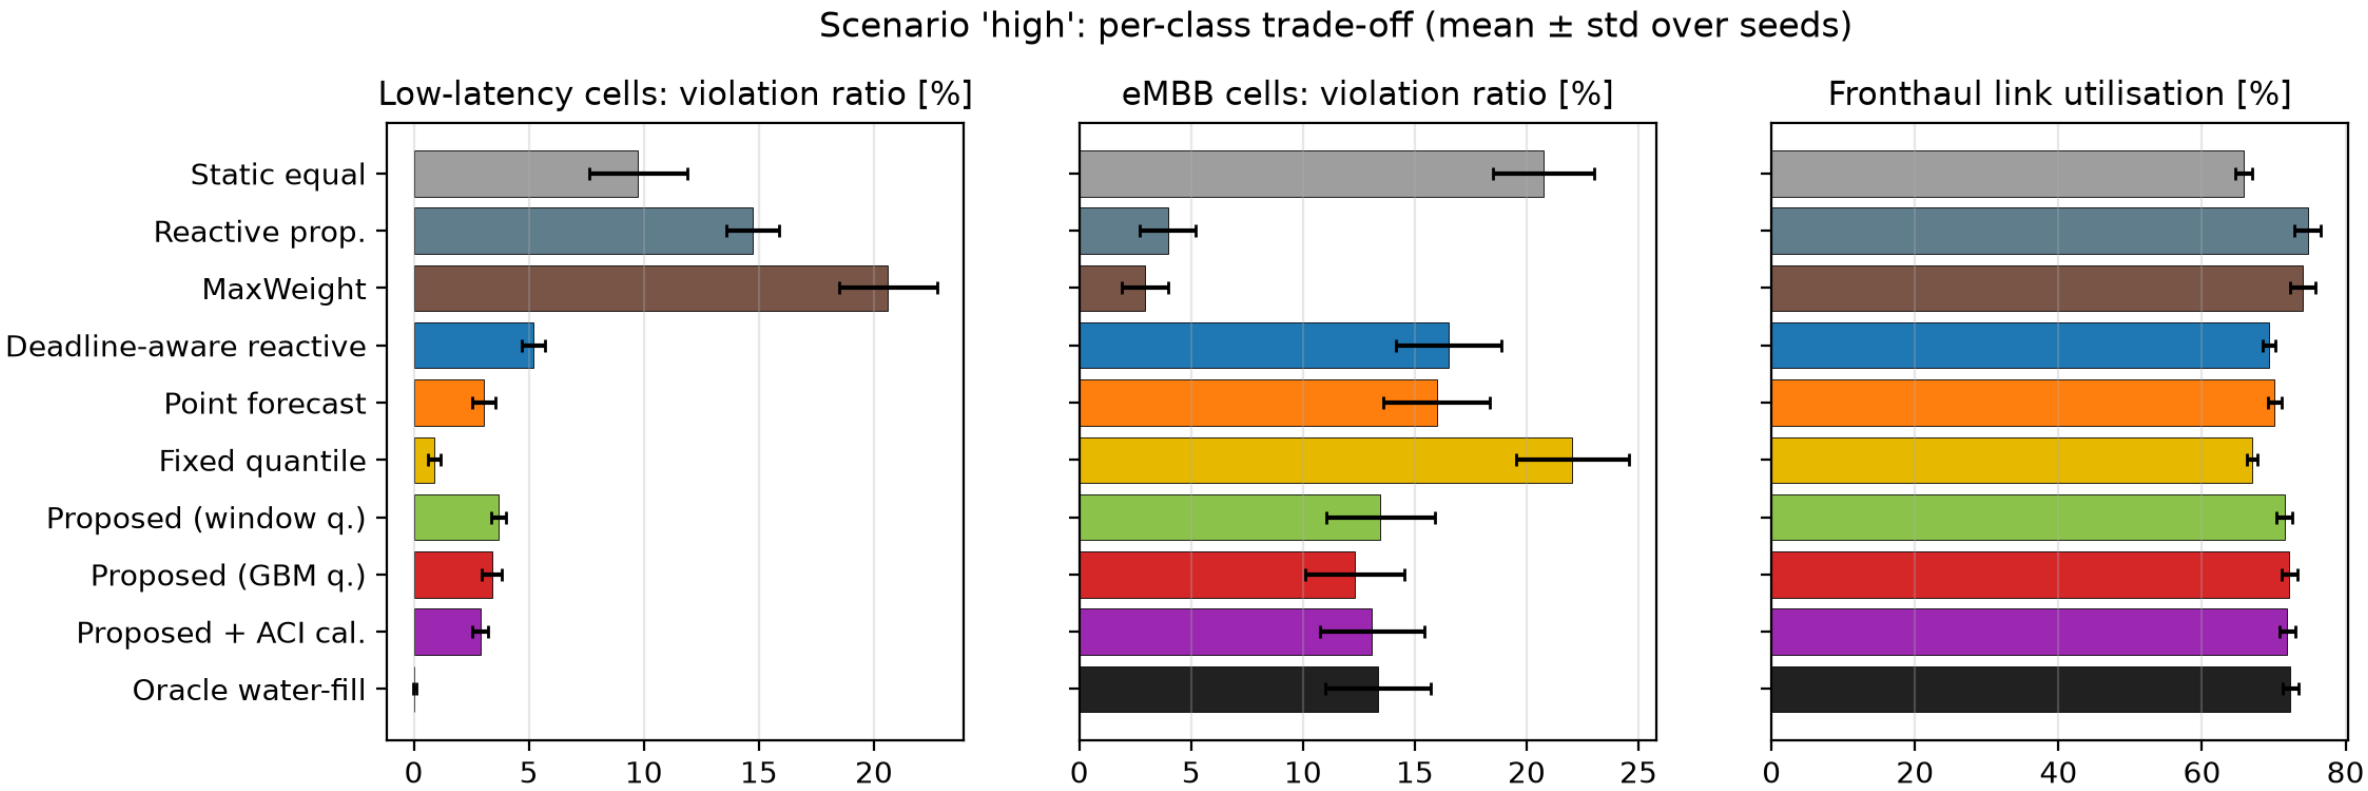





- **Wins 20/20 seeds vs point forecast, deadline-aware reactive, MaxWeight and fixed quantile in every scenario (Holm-corrected Wilcoxon $p < 0.01$)**
- vs reactive proportional: 20/20 in four scenarios; 16/20 under flash crowds (−4 %, $p = 0.03$)
- **ML ablation (same rule, simple window quantiles instead of the trained model):**
  - the model adds 2-37 %; least under the held-out shift (2 %) → the rule, not the model, carries most of the gain
- Online calibration: no gain on normal traffic, harmful under shift (it over-reserves for low-latency cells and starves eMBB) → negative result
- 20 seeds: paired per-seed differences, 95 % bootstrap CI, two-sided Wilcoxon signed-rank test with Holm correction (full table in the paper)

# Where the gain comes from: low-latency vs eMBB (high load)

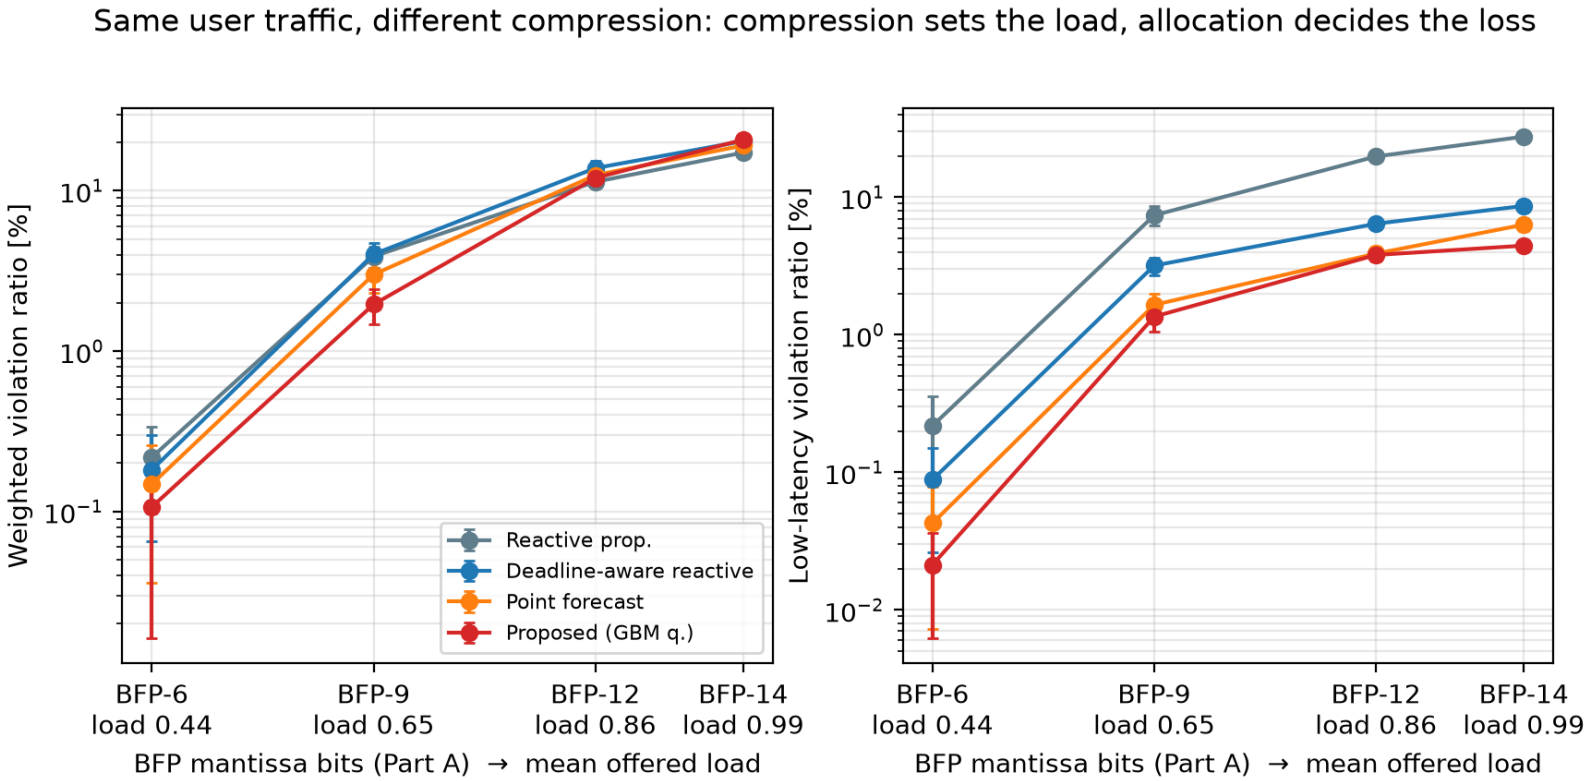





- Reactive proportional: low-latency loss 14.7% vs 3.4% for proposed — controllers that ignore deadlines sacrifice the low-latency cells
- Among deadline-aware controllers, proposed has the lowest eMBB loss (12.3% vs 16.0% for point forecast) and the highest utilisation (72.2%)
- Even the oracle loses 13.4% of eMBB bits at 0.80 load: budgets fixed for 10 ms cannot follow heavy bursts → the control interval matters

# Both levers on the same traffic: compression × allocation

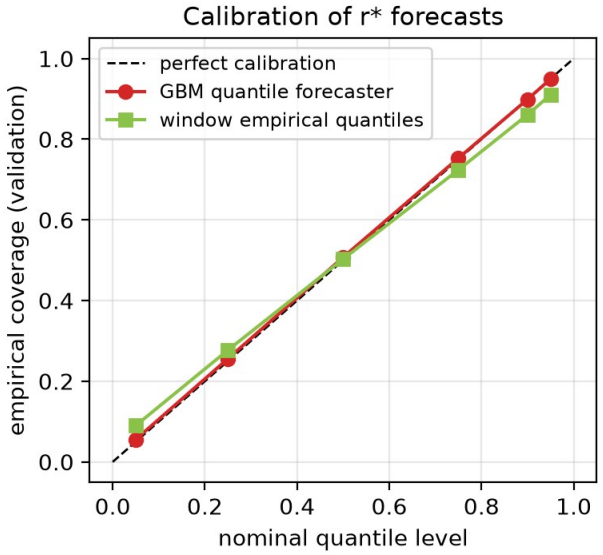
A + B

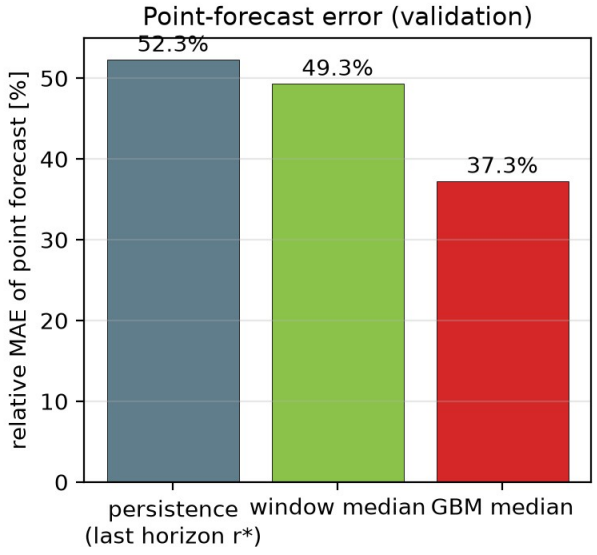


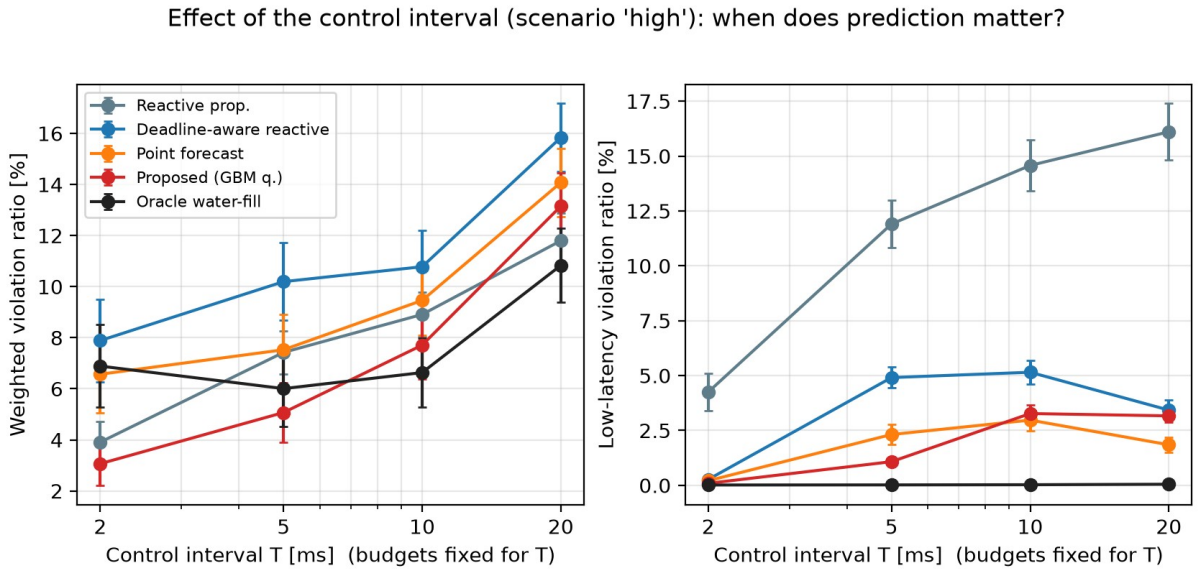

- **Same user traffic, 10 seeds, scenario 'moderate'; only the BFP width changes**
  - BFP-6: load 0.44 → proposed 0.11%, reactive prop. 0.22%; best = proposed
  - BFP-9: load 0.65 → proposed 1.96%, reactive prop. 3.90%; best = proposed
  - BFP-12: load 0.86 → proposed 12.01%, reactive prop. 11.37%; best = reactive prop.
  - BFP-14: load 0.99 → proposed 20.81%, reactive prop. 17.29%; best = reactive prop.
- **Compression sets the regime; allocation helps most between ~0.5 and ~0.8 load**
- At ~1.0 load no allocation rule helps — compress more (Part A) or add capacity
- Not modelled: the signal-quality cost of fewer bits (Part A's NMSE curves cover that side)

# Ablations: forecast quality and control interval






- Validation: the 5 / 50 / 95 % quantiles are exceeded 5.6% / 50.7% / 94.9% of the time (well calibrated); median error 37% vs 52% for persistence
- **Interval sweep (scenario high, 10 seeds). Reduction by proposed vs deadline-aware reactive / vs reactive proportional / vs oracle water-fill:**
  - T=2 ms: +61% / +22% / +55% | T=5 ms: +50% / +32% / +16% | T=10 ms: +28% / +13% / -16% | T=20 ms: +17% / -11% / -21%
- At T = 2 ms the rule even beats the oracle water-fill (perfect single-number forecast) → the allocation rule, not only the forecast, does the work
- At T = 20 ms (5× the low-latency deadline) reserving the low-latency peak rate wastes capacity and reactive proportional wins → the method is for intervals up to a few times the shortest deadline

# Limitations and honest status



- **Part A: compression numbers are quoted from the earlier report and saved figures; the src/ package is not in the repo, so they are not re-run. The matrix-inversion benchmark is re-runnable (needs PyTorch)**
- Part A: benchmark matrices are random and well-conditioned; real channel matrices can be worse
- **Part B: synthetic traffic only — no real radio-unit traces, so no real-world validation is claimed**
- Not a universal win: only ≈ 4 % better than reactive proportional under flash crowds; loses to it at 20 ms intervals, at ~0.85+ load and when telemetry is ≥ 5 ms old; the ML model adds ≈ 2 % under the held-out shift
- Online calibration (ACI) did not help — reported as a negative result; the simple-quantile variant is the current fallback
- 20 test seeds, paired bootstrap CIs and Holm-corrected Wilcoxon tests; still one synthetic traffic generator, so no real-world claim
- Simplified link: a budget-enforced pipe, no switch queue, one traffic class per cell, no HARQ feedback; 'violation' is a DU queueing deadline, not end-to-end latency
- Decision time ≈ 3.2 ms per 10-ms epoch in Python (reactive baselines ≈ 0.1 ms); a real DU would need a compiled version
- Proofs cover feasibility, equal weighted risk, backlog-first and bounded loss under forecast miscalibration — no stability or closed-loop optimality claims

# Status and remaining work toward the final evaluation



## Done and checked

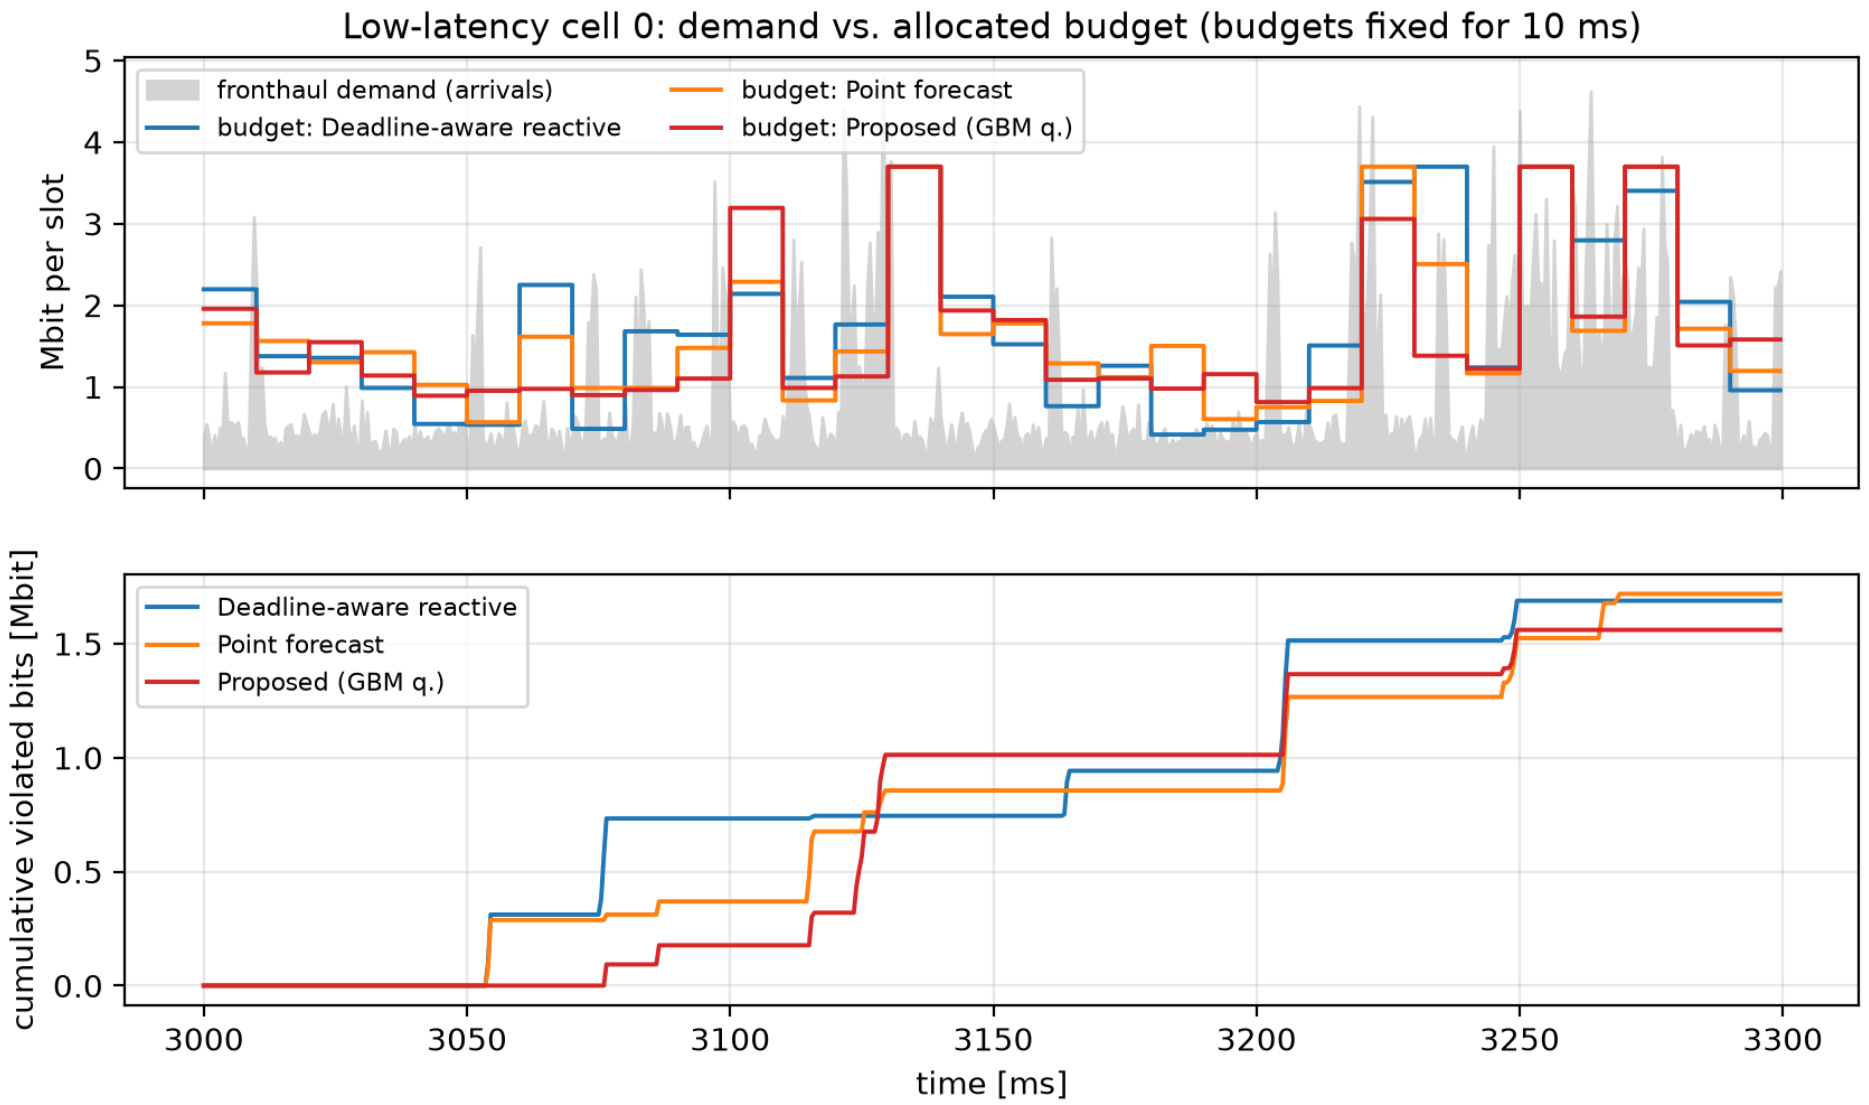
- Part A: compression study (BFP, SVD, RAS-BFP, CSEE, ACAFS) — 4 figure suites + report
- Part A: matrix-inversion benchmark (LU, QR, SVD, Newton-Schulz, learned inverters; n = 10/100/500) — CSV + figure
- Part B: slot-level shared-link simulator with deadlines, priorities, telemetry delay (14 tests)
- Part B: 5 reference controllers, tuned on validation seeds; quantile forecaster, calibrated on validation
- Part B: KKT allocation rule with feasibility proof; 200-run comparison, seed-by-seed analysis, interval sweep
- A + B: compression × allocation experiment; all figures and this deck regenerated by scripts

## Planned (not done yet)

- Part A: restore the compression src/ package into the repo and re-run Exp 1-4; benchmark on realistic channel matrices
- A + B: make the compression width a second knob of the allocation rule (trade signal quality for load per cell)
- Part B: real or public traffic traces → re-check the forecaster
- Part B: skill-weighted mix of trained and simple quantiles as the fallback (instead of ACI)
- Part B: switch-level packet queue and timing window instead of a budget-enforced pipe; more seeds
- Compiled (C++/numba) decision so the 3 ms Python time fits a real DU

# Backup: one trace — budgets vs demand of a low-latency cell





# Backup: code layout (Part B)



| | |
|---|---|
| **config.py** | units, 7-2x fronthaul constants, cells, link |
| **demand.py** | deadline-feasible rate r*, backlog rate |
| **controllers.py** | 5 reference controllers, water-fill |
| **proposed.py** | KKT rule + optional calibration |
| **train.py** | train on training seeds, check on validation |
| **make_figures.py, make_slides.py** | all figures and this deck from CSV |
| **traffic.py** | synthetic arrivals: smooth + bursts + slow changes |
| **simulator.py** | slot loop, queues with ages, drops, metrics |
| **forecast.py** | features, quantile boosting, fast tree evaluation |
| **scenarios.py** | train / validation / test seeds, 5 scenarios |
| **run_experiments.py** | tune / main / sweep / compression / demo |
| **matinv_bench/ (Part A)** | matrix-inversion benchmark: classical + learned |

- Flow: traffic → simulator ↔ controller (observation 2 slots old → budgets) → metrics / raw JSON → figures → slides
- ML flow: training seeds → features / targets → 6 boosted models (~12 s CPU) → validation check → frozen model → test seeds
- Python 3 with numpy / pandas / scikit-learn / matplotlib; full reproduction ≈ 10 min on a 4-core laptop

# Backup: main score [%], mean ± std over 20 test seeds



| Method | Low 0.50 | Moderate 0.65 | High 0.80 | Flash crowds | Shift (held-out) |
|---|---|---|---|---|---|
| Static equal share | 4.45 ± 1.05 | 9.65 ± 1.36 | 15.46 ± 1.51 | 19.15 ± 2.34 | 22.31 ± 2.99 |
| Reactive proportional | 0.83 ± 0.27 | 4.11 ± 0.69 | 9.20 ± 0.99 | 12.27 ± 2.23 | 11.70 ± 2.24 |
| MaxWeight (queue-weighted) | 0.90 ± 0.28 | 4.85 ± 0.79 | 11.54 ± 1.27 | 15.24 ± 3.02 | 12.31 ± 2.44 |
| Deadline-aware reactive | 0.52 ± 0.21 | 4.24 ± 0.78 | 11.03 ± 1.33 | 14.51 ± 3.06 | 12.56 ± 2.81 |
| Point forecast + water-fill | 0.27 ± 0.17 | 3.23 ± 0.73 | 9.70 ± 1.34 | 13.33 ± 3.09 | 11.30 ± 2.81 |
| Fixed global quantile + water-fill | 0.34 ± 0.23 | 4.35 ± 1.02 | 11.76 ± 1.44 | 15.09 ± 3.17 | 12.85 ± 3.32 |
| Proposed rule, simple quantiles (no ML) | 0.25 ± 0.13 | 2.57 ± 0.64 | 8.72 ± 1.34 | 12.67 ± 3.19 | 10.22 ± 2.88 |
| Proposed (GBM quantiles + KKT rule) | 0.15 ± 0.13 | 2.15 ± 0.60 | 8.00 ± 1.27 | 11.76 ± 3.08 | 10.04 ± 2.71 |
| Proposed + online calibration | 0.17 ± 0.13 | 2.17 ± 0.61 | 8.14 ± 1.30 | 12.27 ± 3.26 | 11.12 ± 3.22 |
| Oracle water-fill (knows the future) | 0.07 ± 0.10 | 1.62 ± 0.56 | 6.89 ± 1.28 | 10.16 ± 2.96 | 6.69 ± 2.60 |

*Source: results/fhlm/main_runs.csv (1000 runs). Seed-by-seed statistics: results/fhlm/main_paired.csv.*

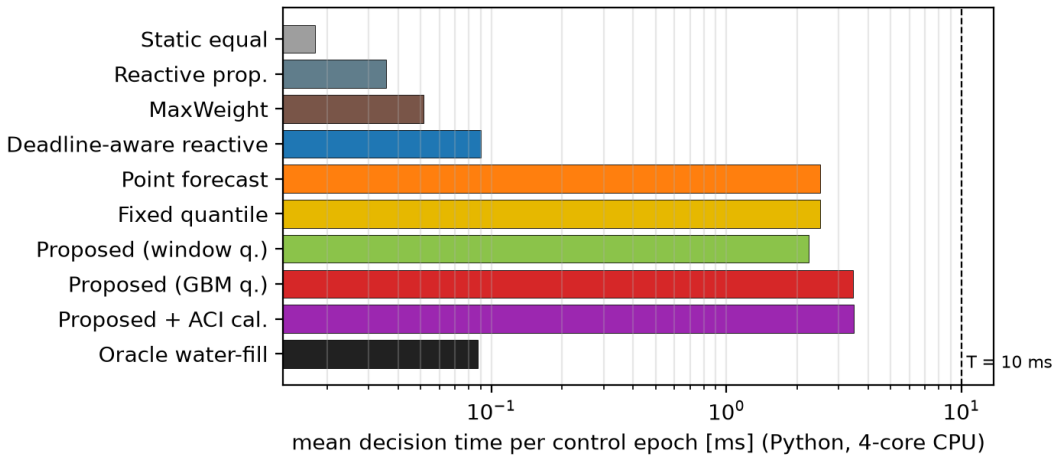
# Backup: where the KKT rule comes from

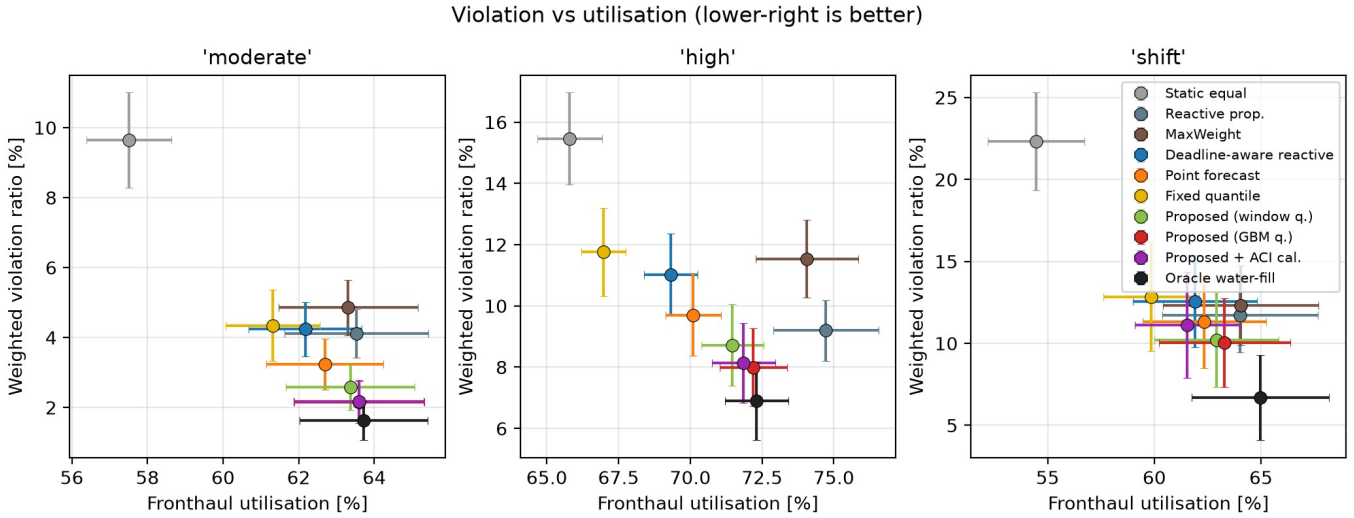


- Objective: $J(r) = \Sigma_i w_i \cdot E[\min(B_i + X_i, r_i)]$, $X_i \sim$ forecast distribution $F_i$ of $r^*_i$, $B_i$ = backlog rate
- Derivative: $d/dr_i\, E[\min(B_i + X_i, r_i)] = P(B_i + X_i > r_i) = 1 - F_i(r_i - B_i) \rightarrow J$ is concave and separable
- Lagrangian with the link constraint ($\lambda$) and the box constraints ($\mu_i$, $\nu_i$)
- Stationarity for an interior $r_i$: $w_i (1 - F_i(r_i - B_i)) = \lambda \Rightarrow r_i = B_i + F_i^{-1}(1 - \lambda / w_i)$
- Clip to $[0, \text{peak}_i]$; $\Sigma_i r_i(\lambda)$ decreases in $\lambda \Rightarrow$ one $\lambda^*$ by bisection (60 steps)
- Property 1 (feasibility): budgets always respect the caps and $\Sigma r_i \leq$ link — checked by test_kkt_allocation_properties
- Property 2 (equal weighted risk): at the optimum every interior cell has $w_i \cdot P(\text{needs one more bit}) = \lambda^*$
- Property 3 (backlog first): $F_i(x) = 0$ for $x < 0 \Rightarrow$ the marginal value below $B_i$ is $w_i \Rightarrow$ backlog is funded before any uncertain demand of a cell with $w \leq \lambda^*$
- A single-number forecast ($F_i$ a step) gives a priority-ordered fill = the point-forecast water-fill baseline
- $F_i^{-1}$: straight lines through the 6 forecast quantiles, extended linearly above 95 %; quantiles sorted to keep them monotone

# Backup: decision time and violation-vs-utilisation trade-off






Decision times are Python on a 4-core CPU inside the simulator (feature construction + 1,800 trees with a vectorised traversal + bisection). The rule itself is < 0.1 ms.

# Backup: references




• O-RAN WG4, Control, User and Synchronization Plane Specification (O-RAN.WG4.CUS.0) [standard] — 7-2x, BFP compression (Annex A.1), T2a windows. 3GPP TR 38.801 [standard]; 3GPP TR 38.901 (TDL-A channel) [standard]; IEEE 802.1CM-2018 [standard].

• L. M. P. Larsen, A. Checko, H. L. Christiansen, 'A Survey of the Functional Splits Proposed for 5G Mobile Crosshaul Networks', IEEE COMST 21(1), 2019. DOI 10.1109/COMST.2018.2868805 [peer-reviewed]

• G. O. Pérez, J. A. Hernández, D. Larrabeiti, 'Fronthaul Network Modeling and Dimensioning Meeting Ultra-Low Latency Requirements for 5G', JOCN 10(6), 2018. DOI 10.1364/JOCN.10.000573; IEEE Access 2019, DOI 10.1109/ACCESS.2019.2923020 [peer-reviewed]

• L. Wang, S. Zhou, 'On the Fronthaul Statistical Multiplexing Gain', IEEE Commun. Lett. 21(5), 2017. DOI 10.1109/LCOMM.2017.2653120 [peer-reviewed]

• S. Lagén, X. Gelabert, A. Hansson, M. Requena, L. Giupponi, 'Fronthaul Compression Control for Shared Fronthaul Access Networks', IEEE Commun. Mag. 2022. DOI 10.1109/MCOM.001.2100959 [peer-reviewed]

• D. Bega, M. Gramaglia, M. Fiore, A. Banchs, X. Costa-Pérez, 'DeepCog', IEEE INFOCOM 2019, DOI 10.1109/INFOCOM.2019.8737488; 'AZTEC', IEEE INFOCOM 2020, DOI 10.1109/INFOCOM41043.2020.9155299 [peer-reviewed]

• F. Kavehmadavani, V.-D. Nguyen, T. X. Vu, S. Chatzinotas, 'Intelligent Traffic Steering in Beyond 5G Open RAN Based on LSTM Traffic Prediction', IEEE TWC 2023. DOI 10.1109/TWC.2023.3254903 [peer-reviewed]

• K. M. Cohen, S. Park, O. Simeone, P. Popovski, S. Shamai, 'Guaranteed Dynamic Scheduling of URLLC Traffic via Conformal Prediction', IEEE WCL 2023 [peer-reviewed]; I. Gibbs, E. Candès, 'Adaptive Conformal Inference Under Distribution Shift', NeurIPS 2021 [peer-reviewed]

• V. Kasuluru, L. Blanco, C. J. Vaca-Rubio, E. Zeydan, 'On the Impact of PRB Load Uncertainty Forecasting for Sustainable Open RAN', arXiv:2407.14400, 2024 [preprint]

• Eckart & Young 1936 (best low-rank approximation = truncated SVD); Halko, Martinsson, Tropp, SIAM Review 2011 (randomised sketching for low-rank factorisation); Willinger et al., IEEE/ACM ToN 1997 (heavy-tailed ON/OFF sources); Hadley & Whitin 1963 (constrained newsvendor); Le Boudec & Thiran, Network Calculus, 2001.